

Trace

A small AI that reads a failed payment and works out what happened to the money — running on the bank's own machines, offline.

Oluwatomiwa Adeaga · deagatech

THE PROBLEM

₦1.07 quadrillion moved. Some of it got stuck.

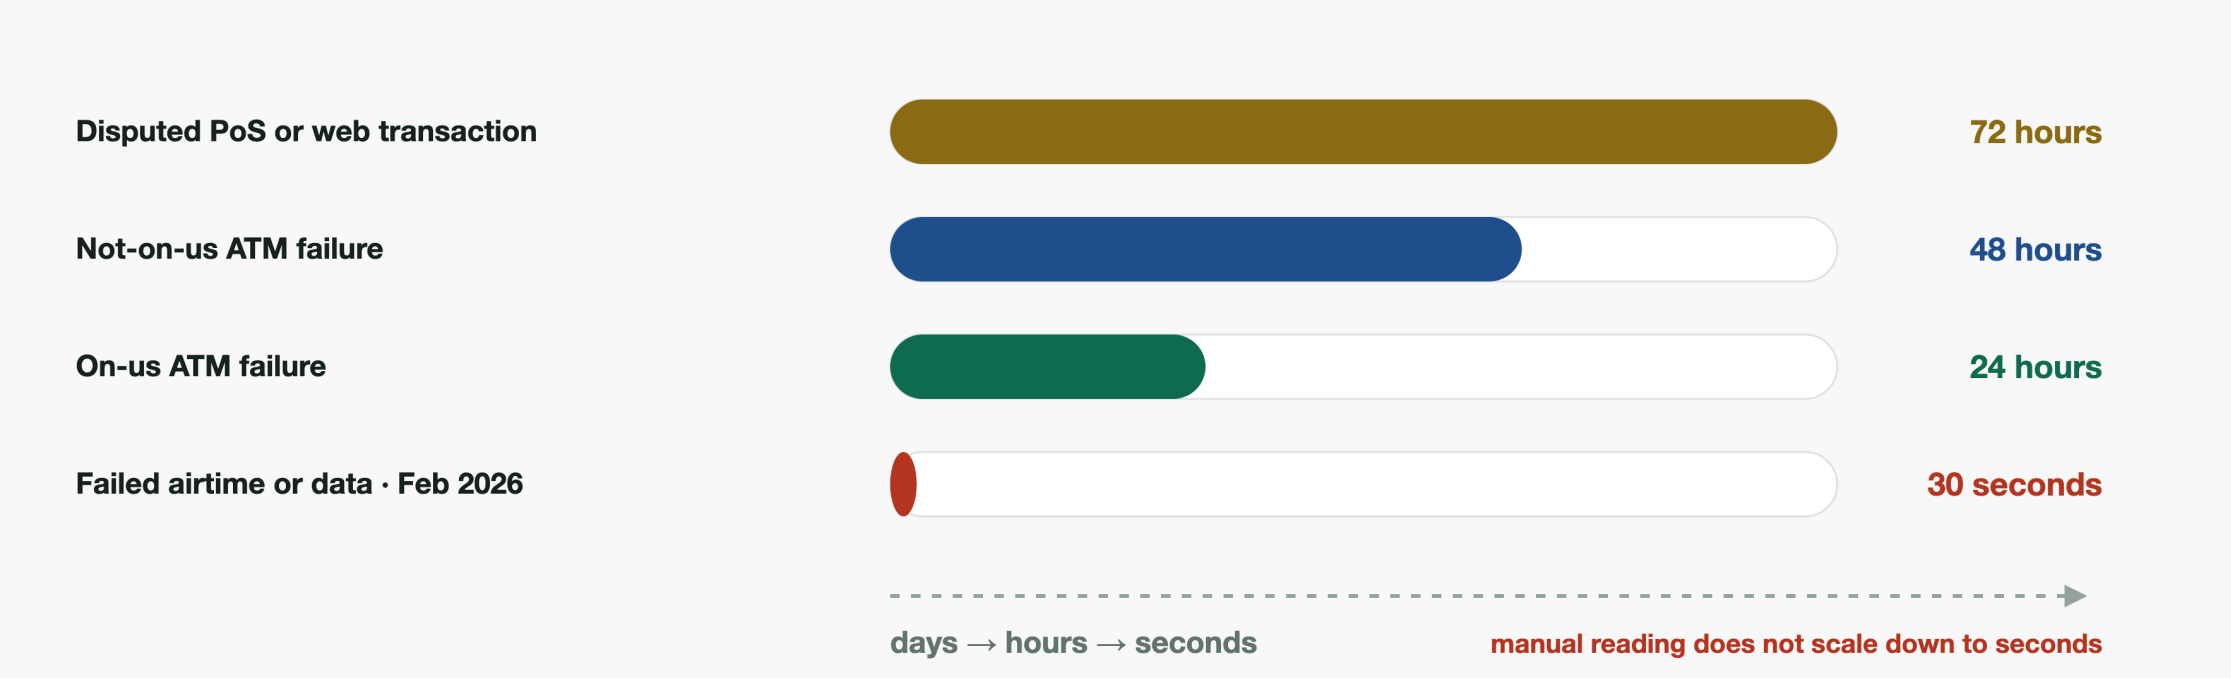


NIBSS 2024 industry figures; CBN consumer protection report, 2025.

WHY IT IS URGENT

And the clock is set by the regulator

These are mandated ceilings, not internal targets.



The direction of travel is one way. In February 2026 the CBN and NCC set thirty seconds for failed airtime.

The answer is never in a database column

EXC-2026-0001930

N402,000.00

switch	91 · Beneficiary bank not available
ledger	DR 402,000.00 · NIBSS/REF957229178847
reversal	raised 20/08 06:15 · RV0765116324 acknowledged: true
settlement	no line for this session

remarks

19/08 17:15 BRANCH: informed cust escalation in progress
19/08 18:15 RECON: cust said she will come to branch tomorrow

what is missing

no credit leg anywhere on the account

The refund was raised.
The switch acknowledged it.

The money never
reached the customer.

A rule that checks
acknowledged == true
closes this case as settled.

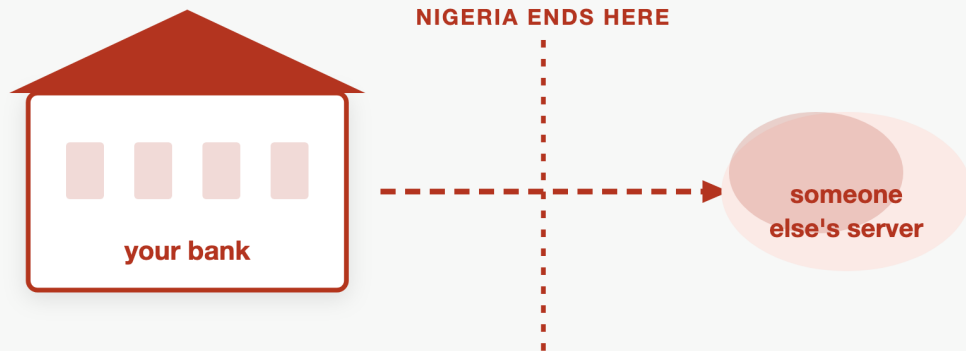
The evidence is an absence.
You have to notice what
is not there.

On this class a rules engine scores 46.4% and an untrained model 9.2%. Trained on the bank's own history: 73.1%.

THE CONSTRAINT EVERYONE HITS

So why not just use ChatGPT?

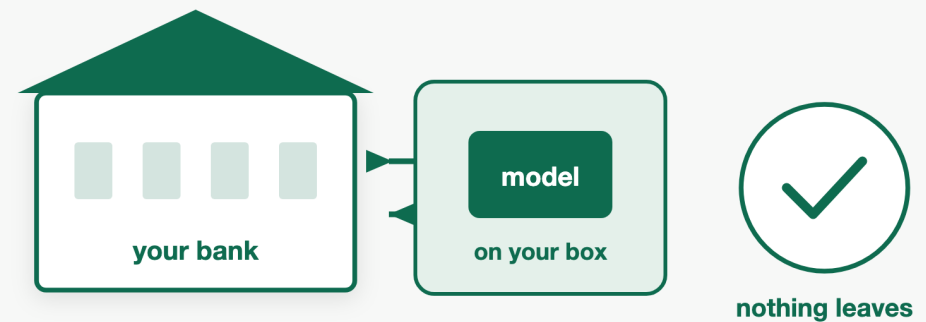
RENTING AN API



Customer names, account numbers and transaction history leave the country.

You own nothing. Price and access can change without you.

RUNNING YOUR OWN



One afternoon of GPU time to train.
Weights, data and test set you keep.

No per-call bill that grows with your own success.

For a Nigerian bank handling customer payment records this is not a preference. It is the difference between a product and a compliance problem.

THE PRODUCT

What Trace is

A case goes in exactly as the desk receives it. An answer comes back that a person can check.

WHAT GOES IN

- the ledger entries
- the switch response
- the settlement file
- what the customer said
- the notes staff typed

One failed payment,
exactly as the desk gets it.



on the bank's own machine
offline



WHAT COMES BACK

- **What went wrong**
one of seven kinds of break
- **Where the money went**
written out, step by step
- **What to do about it**
reverse, credit, recall, close, escalate
- **How sure it is**
so the doubtful ones can be found
- **Whether a person must see it**
and it can always say yes

An answer a person can check, not just a label.

The evidence is an absence

It reads the free text and the ledger together, so it can notice what is not there.

The data cannot leave

It is small enough to run on the bank's own hardware. Nothing is sent anywhere.

The regulator's clock

It answers in about a second, so the deadline stops being the binding constraint.

THE RESULT

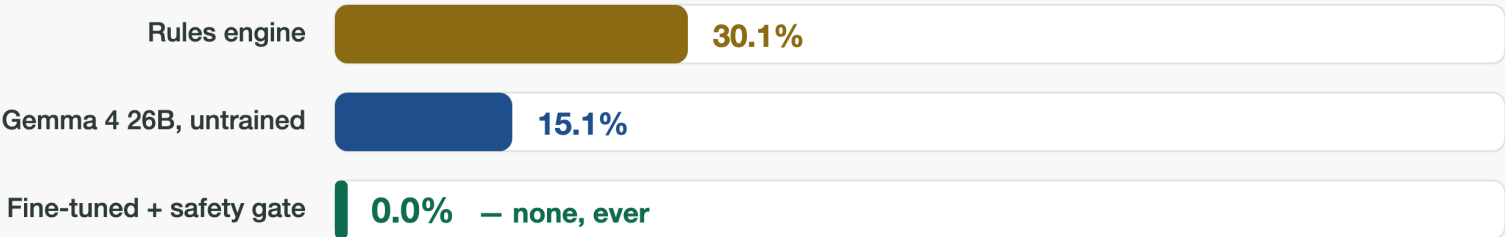
A 2B model beats a 26B model that never saw your data

GETTING THE PROBLEM RIGHT



FRAUD CLOSED WITH NOBODY LOOKING

lower is better

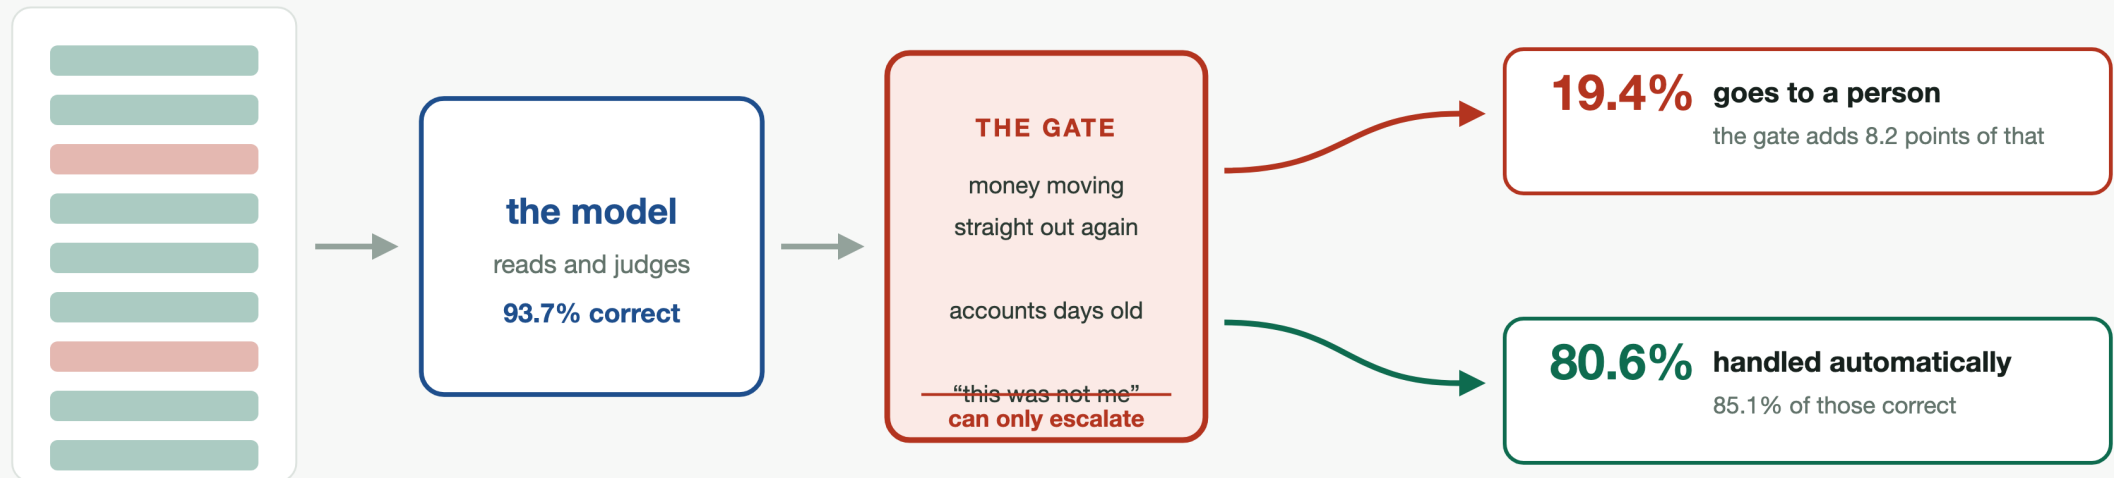


1,000 held-out cases, none seen during training. Every prediction published at github.com/adg-all-day/payment-exceptions-triage

THE SAFETY GATE

It can never quietly close a fraud

EVERY CASE



The model may be wrong. The system may not be dangerous.

Fraud closed with nobody looking: 30.1% under a rules engine, 0.0% here.

How Trace makes money

Priced per institution, not per transaction — and the reason why is the same reason the product exists.

TWO LINES OF REVENUE

Build

paid once
per institution

One engagement per bank. We train the model on their own closed exceptions and hand over the weights, the model and the test set that proves it works.

Keep

paid annually
per institution

Fraud patterns move and detection rules change, so the model is retrained each year against a refreshed test set, with support.

Why not charge per transaction?

Because nothing leaves the building, we cannot meter what we cannot see.

WHO PAYS

44

deposit-taking banks
the direct market

916+

microfinance banks
same problem, no team to build it

then

insurers, PSPs, any regulated back office
same shape of work

Nobody can serve them all with one model.

Our own numbers prove it: the generic model scores 69.5%.

Why me

One person, and an unusual pair of jobs.

I build the tests that measure AI

At Turing and AfterQuery I design evaluation datasets and scoring rubrics for frontier AI labs — task formats, edge cases, acceptance criteria. That is exactly what this project is: a benchmark, built before the model, so the score could not flatter it.

I have worked on fintech projects

As a mobile developer, on products people use to move money. Enough to know what a failed transfer looks like from the customer's side, and what it looks like on the ledger.

Most people entering have built an application. I build the measurement that says whether a model works — and I have worked on the rails this one sits behind.

DEAGATECH

The ask

Not prize money.

Anonymised history

Six months of closed exceptions from one break type.
Account numbers hashed, names removed — the text is what matters.

Nothing identifying ever leaves your building.

One person, six weeks

A named contact in operations who can tell us when an answer is wrong and why.

This is the part that cannot be bought.

A number to beat

What your desk resolves per analyst per day today, and how often a case is reopened.

We will publish whether we beat it or not.

You keep the model, the weights, and the test set. If you never speak to us again, it still runs.