


Payments exceptions model — 1. Where the training data comes from

Code decides the ANSWER so the label is certain. Agents write the PROSE so the evidence cannot be pattern-matched. Neither alone is enough.




Three rules learned the hard way, each after a measured failure:

1. Evidence must be PRESENT — deleting the generator's remarks left 47% of fraud cases unanswerable.
2. Evidence must be in NATURAL language — keeping the template put the rules baseline back to 96.2%.
3. The answer may only cite evidence the case CONTAINS — 42% of fraud answers named onward hops that were not there.

Payments exceptions model — 2. Training, measuring, delivering

One held-out set of 1000 cases scores four systems. The fraud auto-close rate is the number that decides whether any of it is sellable.



TRAIN (4000) reaches the TRAINER only. TEST (1000) reaches the SCORER only. Nothing crosses.
The test split never receives label noise — it is the ruler, and a bent ruler measures nothing.